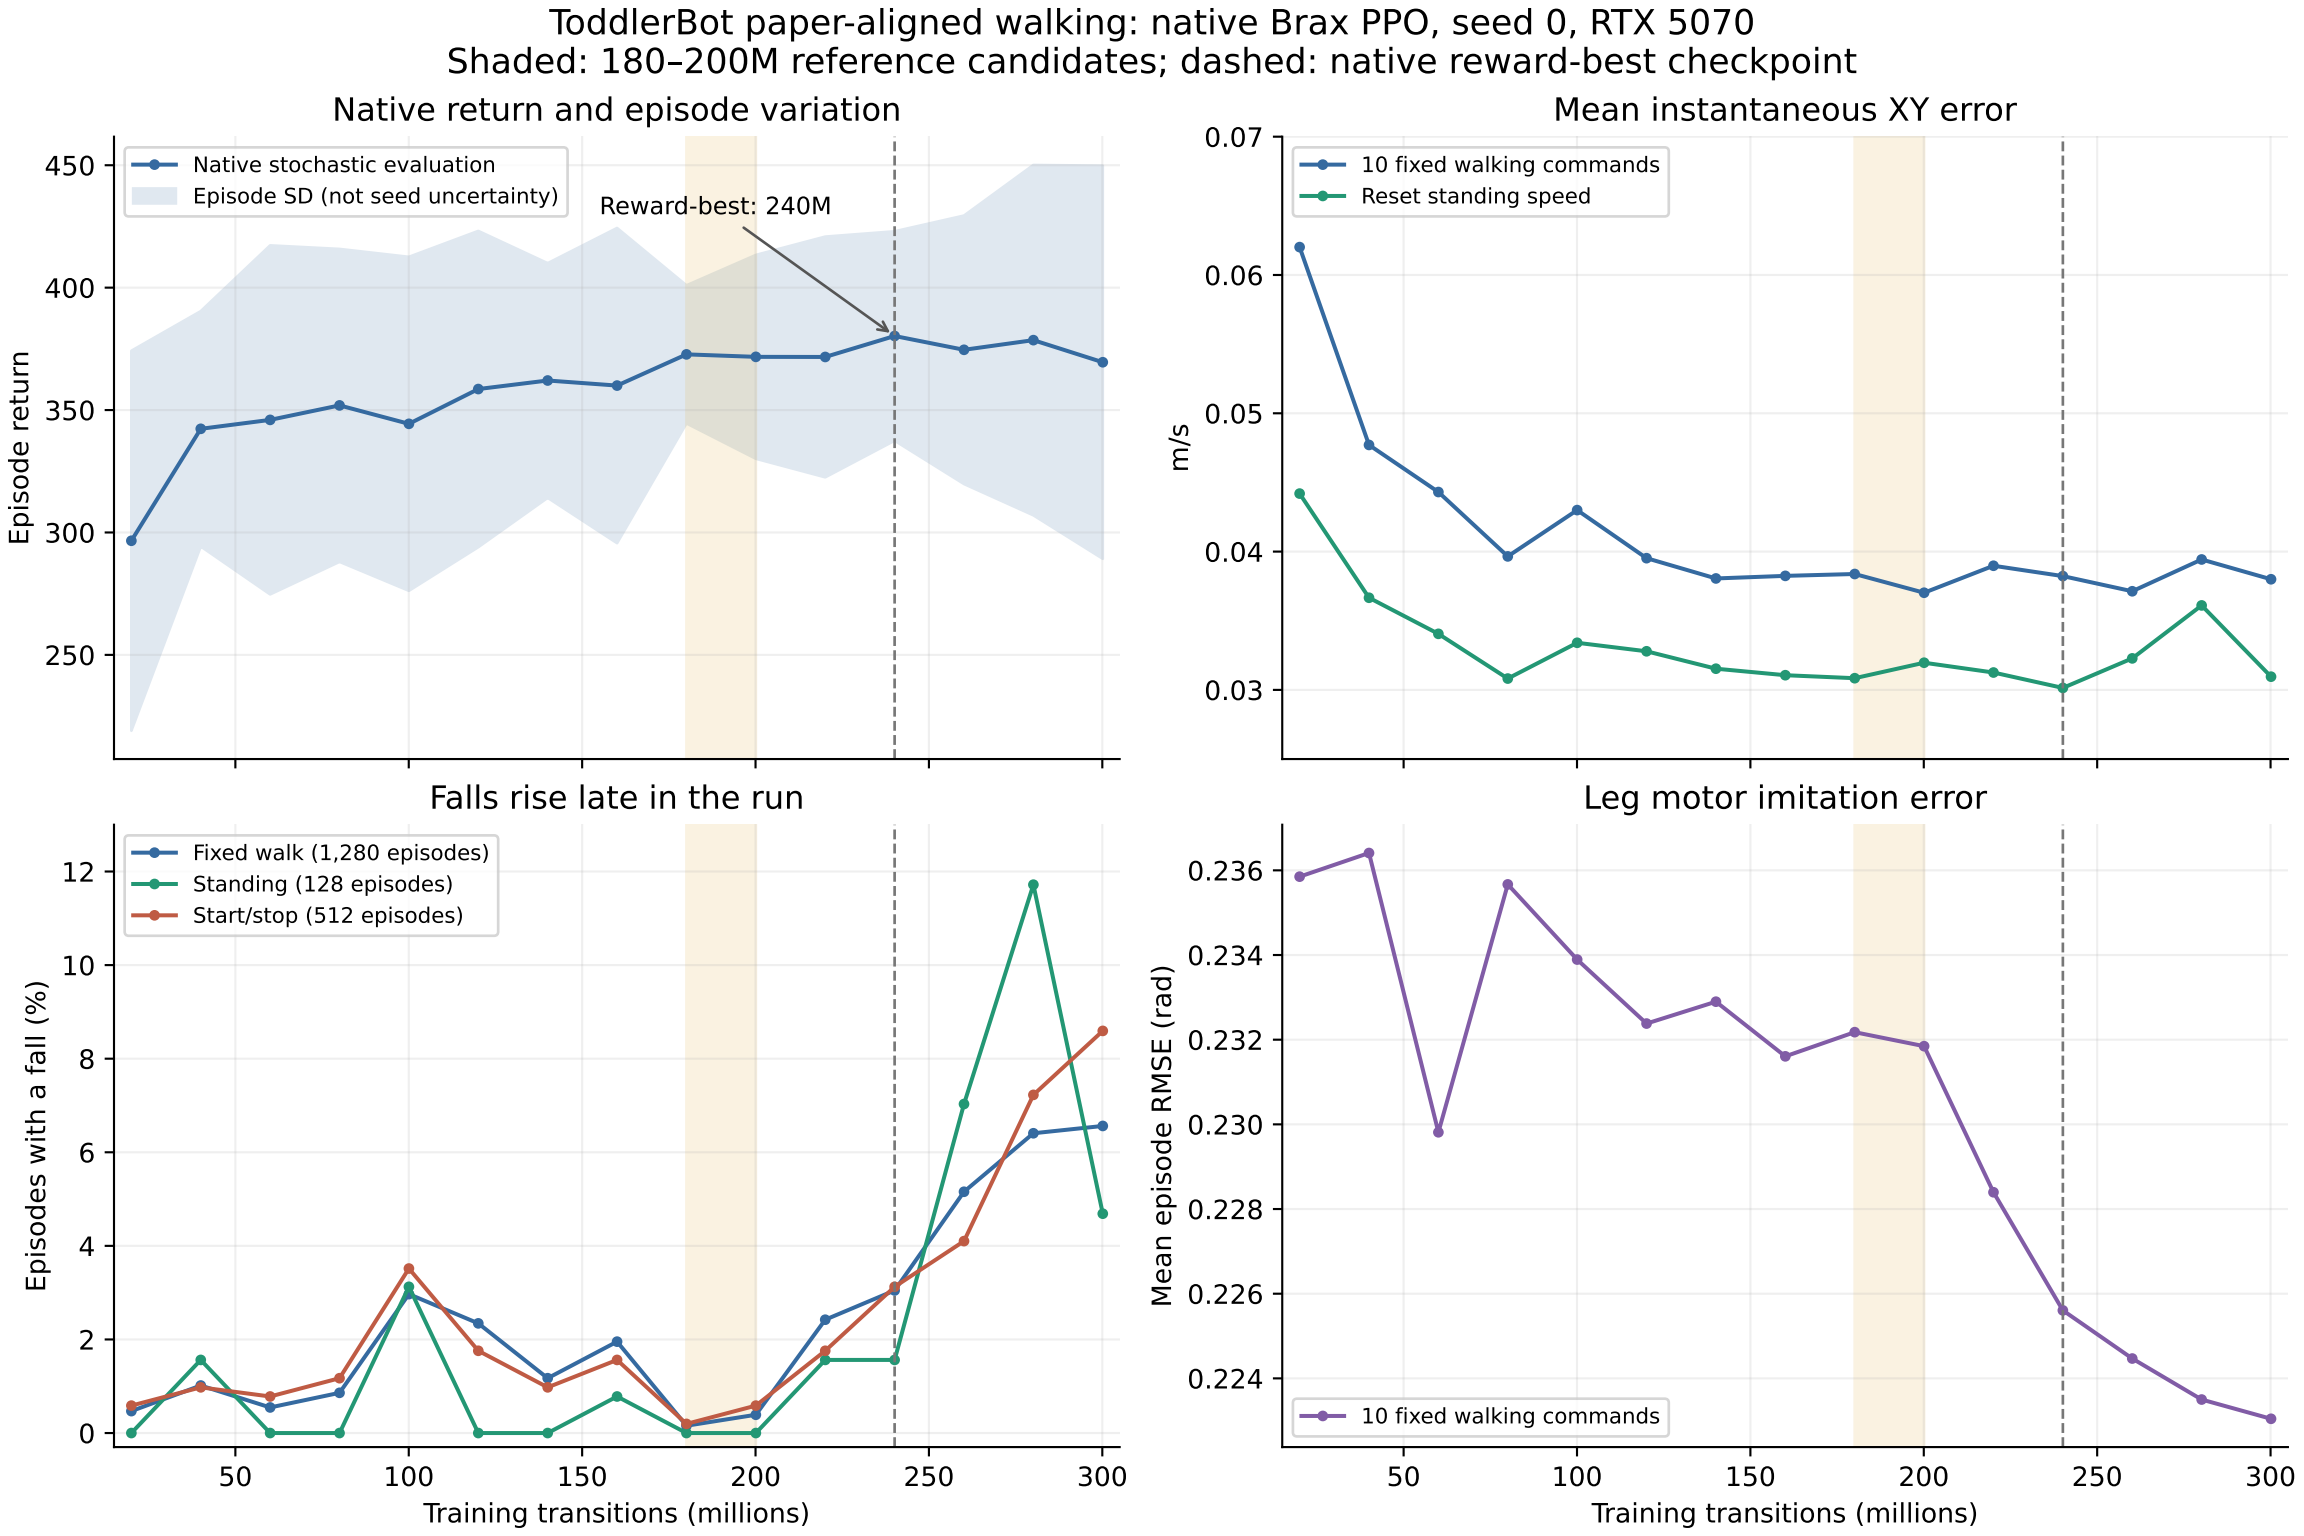
ToddlerBot paper-aligned walking: native Brax PPO, seed 0, RTX 5070
Shaded: 180–200M reference candidates; dashed: native reward-best checkpoint
Native return and episode variation
Native stochastic evaluation
Episode SD (not seed uncertainty)
Reward-best: 240M
Episode return
Mean instantaneous XY error
10 fixed walking commands
Reset standing speed
m/s
Falls rise late in the run
Fixed walk (1,280 episodes)
Standing (128 episodes)
Start/stop (512 episodes)
Episodes with a fall (%)
Training transitions (millions)
Leg motor imitation error
Mean episode RMSE (rad)
10 fixed walking commands
Training transitions (millions)